

Self-Dual BKZ on NTRU: Dimension-Onset, Fatigue, and a Cross-Engine Reference-Free DSD Observable

Brendan Chambers-Bourgeois 

Independent Researcher, Melbourne, Australia

Abstract

This technical report extends the self-dual BKZ (SD-BKZ) reduction benchmark of the companion paper from LWE-Kannan lattices to NTRU , and from a single enumeration engine (fplll) to a second, sieving engine (G6K). We report three empirical results. First, on circulant NTRU lattices the SD-BKZ advantage exhibits a sharp *dimension-onset*: tied with BKZ for small n , a high-variance spike around $n \approx 71\text{--}73$ (mean advantage up to $+9.7$ nats), then a tight $+1.5$ -nat plateau for $n \geq 79$. Second, the defensible NTRU observable is *reference-free*: signed $d(\text{LN})$ is a reference artifact on NTRU (no canonical fixed point), but the BKZ-vs-SD-BKZ profile divergence / dense-sublattice-discovery (DSD) onset is robust, and the SD-vs-BKZ DSD-onset gap grows with dimension ($0\% \rightarrow 27\%$ over $n = 67\text{--}113$) as q crosses the fatigue point. Third, the cross-engine study confirms the companion paper’s thesis from the other side—the Root Hermite Factor is blind to the SD-BKZ difference while $d(\text{LN})/\text{DSD}$ sees it—now on a sieve oracle, and finds the effect to be regime-dependent. We also re-validate, on a new modulus family, an fplll GSO numerical fix. All experiments are reproducible via a containerised, determinism-gated benchmark.

Keywords: lattice reduction, BKZ, SD-BKZ, NTRU , overstretched, fatigue, dense sublattice discovery, Rankin profile, sieving, G6K

1 Introduction

The companion paper [CB26] showed, on LWE-Kannan lattices, that the Root Hermite Factor (RHF) is structurally blind to differences between BKZ and self-dual BKZ [MW16] that the Li–Nguyen Rankin profile distance $d(\text{LN})$ [LN20] makes visible. This report asks two follow-up questions a data paper is well suited to answer: (i) *does the SD-BKZ phenomenology generalise beyond LWE to NTRU?* and (ii) *does it reproduce on a fundamentally different reduction oracle—sieving (G6K) rather than enumeration?*

Our contributions are primarily empirical and infrastructural:

- A characterisation of the SD-BKZ advantage on circulant NTRU lattices [DvW21], including a sharp dimension-onset transition (Section 4).
- Evidence that the appropriate NTRU observable is reference-free, and a *dense-sublattice-discovery (DSD) onset* whose SD-vs-BKZ gap grows with dimension as q approaches the fatigue point (Section 5).

Email: brendanchambersbou@gmail.com (Brendan Chambers-Bourgeois)



- A cross-engine study (Section 6): G6K wired as a second engine behind a reproducibility-gated seam, a self-dual pump-and-jump construction validated against fpIII’s SD-BKZ, and confirmation that RHF-blindness is engine-independent but regime-dependent—consistent with independent observations of RHF-equality under a non-exact oracle [Row24].
- An independent re-validation of an fpIII GSO catastrophic-cancellation fix on a new (NTRU) modulus family (Section 7).

This is a *technical report / data paper*: the emphasis is on the experimental apparatus, the data, and reproducibility, rather than on new theory. Closest prior work is Deng and Jia’s self-dual BKZ over a sieve oracle [J⁺23]; we differ in measuring the SD-vs-BKZ structural difference via the reference-free profile observable rather than swapping the oracle for speed. White [Whi26] studies dynamic block-size strategies within the same Li–Nguyen framework, complementary to the fixed- β , q -swept regime characterised here.

2 Background

2.1 Circulant NTRU lattices and fatigue

We use the circulant NTRU construction of Ducas and van Woerden [DvW21]. For a prime n and modulus q , sample ternary secrets $f, g \in \mathbb{Z}_q[x]/(x^n - 1)$ (uniform $\{-1, 0, 1\}$, coordinate variance $\sigma^2 = 2/3$) with f invertible, set $h = g/f \bmod q$, and form the $2n \times 2n$ basis

$$B = \begin{pmatrix} qI_n & 0 \\ H & I_n \end{pmatrix},$$

where H is the circulant matrix of h . The secret (f, g) is an unusually short vector in the row span; recovering it breaks the key. Two structures distinguish B from a random q -ary lattice: a q -vector *head* (rows of qI_n that survive reduction as length- q vectors) and a *dense secret sublattice* spanned by the rotations of (f, g) . A BKZ-reduced NTRU profile therefore follows the *ZGSA* (Z-shaped) model—a flat head at $\log q$, then a GSA slope—rather than the single-slope GSA line of a random q -ary lattice.

As q grows at fixed n , the secret sublattice becomes relatively denser and easier to detect: at the *fatigue point* $q_{\text{fat}}(n) \approx 0.004 n^{2.484}$ [DvW21] the attack transitions from the standard regime (secret found only through full SVP) to the *overstretched* regime, where reduction discovers the dense sublattice (*dense-sublattice discovery*, DSD) well before solving SVP. The geometry of this transition—and how self-dual reduction moves through it—is the object of this report.

2.2 BKZ, SD-BKZ, and the two engines

Block Korkine–Zolotarev (BKZ) reduction repeatedly calls an SVP oracle on length- β projected sublattice windows, sweeping left to right. Self-dual BKZ (SD-BKZ) [MW16] alternates a forward (primal) pass with a backward pass on the reversed dual basis, so each tour reduces the basis from both ends. On LWE–Kannan lattices the companion paper [CB26] found the two produce bases of indistinguishable Root Hermite Factor yet measurably different Rankin profiles; this report asks how that difference behaves on NTRU and across two oracle implementations.

We use two engines. **fpIII** supplies an exact enumeration SVP oracle; its self-dual variant (`SD_VARIANT`) is our ground truth. **G6K** [ADH⁺19] supplies an approximate *sieving* oracle through the pump-and-jump BKZ tour (`pump_n_jump_bkz_tour`); we realise SD-BKZ on it as a primal pump-and-jump pass followed by a dual pass run under G6K’s `dual_mode` (operations on the dual basis, bounds reflected about `full_n/2`). The two

oracles differ fundamentally—deterministic enumeration versus randomised sieving—which is what makes a phenomenon that survives both engines credible. Sieving carries a fixed setup cost that makes small β uninformative, so all cross-engine experiments use $\beta = 40$, the smallest block size at which the sieve does non-trivial work; this is also the regime in which related q -ary forgery attacks operate [DEP23].

2.3 The reference-free DSD observable

The companion paper’s metric is $d(\text{LN}) = \frac{1}{m} \sum_i |r_i - \rho_i|$, where r_i is the slope-removed Gram-Schmidt log-norm profile and ρ_i is the Li-Nguyen BKZ fixed point [LN20]—a GSA line valid for a *random* q -ary lattice. NTRU is not random q -ary (Section 2.1): its reduced profile follows the ZGSA Z-shape, for which no closed-form Li-Nguyen fixed point exists. Comparing an NTRU profile to the LWE line therefore conflates structure with reduction dynamics, and the *signed* advantage $d(\text{LN})_{\text{BKZ}} - d(\text{LN})_{\text{SD}}$ can flip sign under a different reference—it is a *reference artifact*.

Two observables survive this. First, the *unsigned* profile divergence $\frac{1}{m} \sum_i |r_i^{\text{BKZ}} - r_i^{\text{SD}}|$ is reference-free (it never names ρ); on NTRU it is large precisely at the transition and reproduces under three independent references. Second, and what we report throughout, is a binary *dense-sublattice-discovery* flag: a reduction has found the secret sublattice when $b_1 = \min_i \log \|\mathbf{b}_i^*\| > 1.5$, where the threshold is the q -independent ternary-secret log-norm floor $\log \sqrt{2n \cdot 2/3} \approx 1.96$ (we also track the count of GS vectors that have collapsed below the q -vector head, $\leq n + 1$). Both are properties of the output basis alone, independent of any modelled reference. The onset q is the smallest modulus at which a given reduction flags DSD; the SD-vs-BKZ onset *gap* is the object of Sections 5 and 6.

3 Experimental Setup

3.1 Generators and engines

NTRU bases are built by `generators/ntru.py` following Section 2.1; every basis is checked structurally at construction ($\dim = 2n$, a qI_n block, and $Hf \equiv g \pmod{q}$) before any reduction. Both engines sit behind a single seam: `run_single` takes a `backend` argument, and `make_backend` returns either an `fpdll` or a G6K backend exposing a common `tour()/gso()` interface. The per-tour reduction loop is then engine-blind—the same loop drives enumeration and sieving, so the two cannot diverge through duplicated control flow. The G6K self-dual tour is a primal `pump_n_jump_bkz_tour` followed by a dual pass executed inside a `temp_params` block that flips `dual_mode`; this was validated against `fpdll`’s `SD_VARIANT` by checking that the SD-BKZ profile-change signature (the GS-slope flattening) agrees in direction and magnitude across engines (`fpdll` -0.00039 vs G6K -0.00040 per tour at $n = 80$, $\beta = 40$).

3.2 Reproducibility and determinism gates

Both engines run in containers built from a digest-pinned base with `snapshot.debian.org` apt pinning and a fixed `-march=x86-64-v3` target, so the binaries are bit-reproducible across machines. Each result is a per-seed JSON carrying the full input parameters and output profile; a SHA-256 manifest gates byte-identity. The `fpdll` path is held to exact-JSON reproducibility by `verify.sh`; the G6K path by `verify_g6k.sh` against a captured reference hash, confirmed on a second machine (the CI runner).

Determinism on the sieve required care (ADRs 005–008). G6K’s multi-threaded sieve has a nondeterministic reduction race, so all sieving runs single-threaded (`threads = 1`); across-seed parallelism still uses all cores, since seeds are independent. The `fpdll` global RNG must be seeded both before basis construction and again before the sieve (G6K

Table 1: SD-BKZ advantage on NTRU: tied $\rightarrow$ onset spike $\rightarrow$ plateau.

n	mean adv	std	seeds > 0
59–67	~ 0	tight	3–9/20
71	+4.35	5.59	16/20
73	+9.68	5.21	19/20
79–89	+1.5	~ 0.2	20/20

samples from it), and the `SieverParams` seed is a no-op that we never rely on. With these, multi-tour G6K runs are bit-identical run-to-run, and—verified at $n = 80$, $\beta = 60$, three tours—per-tour re-seeding is itself a no-op: the construction seed alone determines all N tours, so the policy is to seed once at construction. The new modulus and engine introduced no schema change to the seed JSON, preserving the companion paper’s SHA baselines.

3.3 Campaigns and parameter grids

Every experiment is a named campaign in `config/sweep.toml`, fixing generator, engine, n -grid, β , q , tour count, and seed count so a single command reproduces a cell. The dimension-onset sweep (Section 4) fixes $q = 97$, $\beta = 20$, 50 tours and sweeps $n \in [59, 89]$ at 20 seeds each. The fatigue and DSD-onset sweeps (Section 5) fix n and sweep q across the fatigue point. The cross-engine sweeps (Section 6) fix $n = 89$, $\beta = 40$, 50 tours and sweep a matched q -grid on *both* engines—the G6K campaign (`ntru_g6k_qsweep`) and the fp111 campaign (`ntru_xeng_qsweep`) write to leaves that share the same (q , precision, tours, n , β) key, so a comparator reads the two engines cell-for-cell. The Kahan re-validation (Section 7) runs $n = 127$ at $q \in \{971, 1087, 1201\}$ through a separate patched-fp111 campaign whose seeds land in their own tree, never overwriting the canonical seeds.

4 The SD-BKZ Advantage Has a Dimension-Onset on NTRU

On circulant NTRU at $q = 97$ (standard regime), $\beta = 20$, 50 tours, the SD-BKZ-vs-BKZ advantage has a *sharp* dimension transition, not a smooth ramp (20 seeds per n):

Three regimes: tied for $n \leq 67$; an onset zone at $n = 71$ –73 where the mean spikes to +4–10 nats with $\sim 25\times$ the plateau variance (SD-BKZ cracks the dense sublattice on most-but-not-all instances); and a stable +1.5-nat plateau for $n \geq 79$ (std ≈ 0.2 , 20/20 wins). The onset spike exceeds the plateau—non-monotonic (Figure 1).

5 Fatigue and the Cross-Dimension DSD-Onset Gap

Before sweeping q we confirm the metric is sound here. The $n \approx 71$ –73 divergence of Section 4 is reference-robust: re-derived from the stored profiles under three independent references (Li–Nguyen, GSA-zero, and the reference-free $\frac{1}{m} \sum_i |r_i^{\text{BKZ}} - r_i^{\text{SD}}|$), it remains large under all three (32.7 nats reference-free at $n = 73$ vs 0.03 at $n = 67$). The *signed* advantage, by contrast, is a reference artifact on NTRU (Section 2.3); accordingly everything below reports the reference-free DSD observable, never the signed quantity.

Sweeping q across the fatigue point at fixed n , the reference-free BKZ-vs-SD-BKZ profile divergence peaks near $q \approx q_{\text{fat}}$. The SD-BKZ DSD-onset occurs at progressively *lower* q than the BKZ onset as n grows—the gap widens with dimension:




Figure 1: NTRU SD-BKZ dimension-onset ($q = 97$, $\beta = 20$, 50 tours, 20 seeds per n , $\pm 1\sigma$ band). The signed advantage is reference-robust at the spike (Section 2.3); shown here for the dimension transition. Recomputed from the per-seed **advantage** field.

Table 2: Reference-free DSD-onset: SD-BKZ reaches dense-sublattice discovery at lower q than BKZ, gap growing with n (~ 15 – 20 seeds/cell).

n	SD onset q	BKZ onset q	gap%
67	146	149	2
79	175	175	0
89	237	281	18
101	426	514	21
113	732	932	27




Figure 2: Reference-free DSD-onset modulus ($b_1 > 1.5$) for SD-BKZ vs BKZ as a function of n ($\beta = 20$), with the SD-vs-BKZ gap% labelled. SD-BKZ reaches dense-sublattice discovery at progressively lower q than BKZ; the gap grows from 0 to 27%. Values mirror Table 2.

$n = 127$ is excluded (fp11 §8 catastrophic cancellation + an off-grid crack). The widening is shown in Figure 2.

6 Cross-Engine: SD-BKZ on a Sieve Oracle

The G6K engine seam and the self-dual pump-and-jump construction are described in Section 3.1; recall that the construction was validated against fp11’s SD_VARIANT by matching the per-tour GS-slope-flattening signature (-0.00039 vs -0.00040). This section reports what the validated sieve engine then shows.

RHF blindness is engine-independent but regime-dependent. On NTRU at $q = 97$, $\beta = 40$, the companion paper’s thesis reproduces on the sieve: in the standard regime the RHF of BKZ and SD-BKZ bases is identical to machine precision while $d(\text{LN})$ differs. The effect is regime-dependent—RHF-blindness breaks only on the rare seeds where SD-BKZ finds a short vector (which RHF reflects). This is consistent with RHF-equality reported under a non-exact (approximate-enumeration) oracle [Row24]; the present work measures the structure RHF misses.

Cross-engine validation of the construction. A dimension-resolved fp11-vs-G6K comparison at $q = 97$ ($n = 89/101/113$, 20 seeds each, $\beta = 40$, 50 tours) confirms the self-dual G6K construction is sound (Table 3): at $n = 101/113$ neither engine fires a short-vector event, and at $n = 89$ the engines overlap (Jaccard 0.50) on the one seed where a short vector genuinely exists (seed 19), where both land the identical first GS log-norm 2.332. The G6K sieve never fires where fp11’s exact oracle does not—there is no spurious

Table 3: Cross-engine soundness at $q = 97$: G6K never fires where fplll does not; the $n = 89$ overlap is the one genuine short-vector seed.

n	G6K fired	fplll fired	Jaccard
89	2/20	1/20	0.50
101	0/20	0/20	1.00
113	0/20	0/20	1.00

Table 4: Matched cross-engine DSD-onset: SD-BKZ reaches DSD before BKZ on *both* engines ($n = 89$, $\beta = 40$, 50 tours, 20 seeds/cell, $b_1 > 1.5$).

q	fplll (enum)		G6K (sieve)	
	SD/BKZ	gap	SD/BKZ	gap
97, 113	0/0	0 (pre-onset)	0/0	0 (pre-onset)
137	8/1	+7	14/4	+10
157	20/18	+2	20/19	+1
≥ 181	20/20	0 (sat.)	20/20	0 (sat.)

sieve firing—so the $n = 89$ short-vector events are real but sparse, not an artifact of the approximate oracle.

The DSD-onset gap reproduces on *both* engines, confound-free. We run the *same* grid ($n = 89$, $\beta = 40$, 50 tours, $q \in \{97, 113, 137, 157, 181, 211\}$, 20 seeds/cell) on both the fplll enumeration engine and the G6K sieve, under the *identical* reference-free DSD criterion: $b_1 = \min_i \|\mathbf{b}_i^*\| > 1.5$ (in log-norm; the q -independent ternary-secret floor $\approx \sqrt{2n \cdot 2/3} \approx 1.96$). On both engines SD-BKZ reaches dense-sublattice discovery at *lower* q than BKZ.

At $q = 137$ SD-BKZ cracks DSD where BKZ does not—8/20 vs 1/20 on fplll, 14/20 vs 4/20 on G6K—the same onset ($q \approx 137$), same direction, same magnitude scale (+7 enum, +10 sieve). This is the DSD-onset gap of Table 2 reproduced across a fundamentally different oracle, and on a matched β : it is therefore *not* a $\beta = 20$ artifact (both engines run $\beta = 40$) and *not* an engine artifact (enumeration and sieving alike). Both onsets sit at lower q than the $\beta = 20$ fplll onset ($q=237$)—stronger reduction lowers the onset.

Crucially the gap is visible *only* reference-free: on the seeds where SD-BKZ cracks DSD, the *signed* $d(\text{LN})$ advantage is large and *negative*—mean -21.9 (fplll) and -30.8 (G6K) at $q = 137$ —because the DSD Z-shaped profile is far from the LWE-derived $d(\text{LN})$ reference. This makes the “signed $d(\text{LN})$ is a reference artifact on NTRU” point concrete and cross-engine: report reference-free DSD, never the signed advantage.

7 Re-Validating the fplll GSO Fix on a New Modulus Family

The companion paper documents a catastrophic-cancellation bug in fplll’s squared-form GSO update at cryptographic moduli, and a Kahan-compensated fix. We re-validate that fix on a new (NTRU) modulus family: on $q \in \{971, 1087, 1201\}$ at $n = 127$, the four contaminated seeds recover from the precision-floor clamp ($b_1 = -345.388$) into the control band (≈ -0.1); the patched build passes all 15 fplll regression tests. The fix is thus validated on two modulus families (the companion paper’s $q = 3329$ LWE, 38% $\rightarrow$ 0% degeneracy, and this NTRU set).

8 Discussion

Three observations bear on how SD-BKZ is modelled in NTRU security estimation. First, the SD-BKZ advantage is *not* a smooth function of dimension: the sharp onset spike at $n \approx 71\text{--}73$ (Section 4), exceeding the plateau, means a security estimate interpolated from a coarse n -grid can miss where self-dual reduction is most effective. Second, the quantity that captures the SD-vs-BKZ difference on NTRU is reference-free: the signed $d(\text{LN})$ advantage flips sign under the reference and must not be reported (Section 2.3); the defensible observables are the unsigned profile divergence and the DSD flag. Third, and most consequential, the DSD-onset *gap*—SD-BKZ reaching dense-sublattice discovery at lower q than BKZ—grows with dimension (Section 5) and reproduces on both an enumeration and a sieving oracle at a matched $\beta = 40$ (Section 6). A model that treats SD-BKZ and BKZ as RHF-equivalent (as the Root Hermite Factor, blind here, would suggest) misses that self-dual reduction crosses the overstretched threshold earlier.

These findings sit alongside two lines of work. The $\beta = 40$ regime and the Z-shaped q -ary profile geometry match the small-modulus forgery setting of Ducas–Espitau–Postlethwaite [DEP23], where the same profile structure is the attack surface. And the move away from an RHF-centric view toward sublattice density echoes recursive lattice reduction [AEP23], which abandons the Hermite-factor framing for a Rankin/dense-sublattice one—the same direction our reference-free observable points.

We are careful about scope. The dimension-onset spike is shown at $\beta = 20$, $q = 97$; the cross-engine gap at $\beta = 40$, $n = 89$. The signed $d(\text{LN})$ is reported nowhere as a result. And $n = 127$ is excluded from the onset trend (Section 5): the fp11l §8 cancellation contaminated those cells, and the circulant crack there is off-grid—it is a validation target (Section 7), not a trend point.

9 Reproducibility

Every result traces to per-seed JSON files under `results/seeds/`, SHA-256-gated for byte-identity and reproducible across machines from digest-pinned containers (Section 3.2). Each experiment is a named campaign in `config/sweep.toml`, so a cell is regenerated by a single `run_campaign` invocation naming the campaign, q , and seed count. The cross-engine soundness records and DSD-onset tables are committed under `results/validation/` with their own reproduction commands; the determinism and self-dual-semantics decisions are recorded in ADRs 005–008. The figures regenerate from the seed data through one figure module each, gated against drift. The companion paper’s SHA baselines are preserved unchanged: the NTRU generator and the G6K engine were added without altering the seed JSON schema.

10 Conclusion

On NTRU, the difference between BKZ and self-dual BKZ that the Root Hermite Factor cannot see is real, reference-free, and consequential: SD-BKZ reaches dense-sublattice discovery at a lower modulus than BKZ, with a gap that grows with dimension and that we observe on two fundamentally different reduction oracles at a matched block size. The companion paper argued the RHF is blind on LWE; this report shows the blindness persists on NTRU and on a sieve, exactly where it matters for the overstretched transition. The apparatus—a determinism-gated, containerised, two-engine benchmark with per-seed evidence—is released so the measurements can be reproduced and extended.

References

- [ADH⁺19] Martin R. Albrecht, Léo Ducas, Gottfried Herold, Elena Kirshanova, Eamonn W. Postlethwaite, and Marc Stevens. The general sieve kernel and new records in lattice reduction. *EUROCRYPT 2019; IACR ePrint 2019/089*, 2019. G6K; pump-and-jump BKZ; dual_mode.
- [AEPsD23] Divesh Aggarwal, Thomas Espitau, Spencer Peters, and Noah Stephens-Davidowitz. Recursive lattice reduction, 2023. arXiv:2311.15064. Abandons rHF for dense-sublattice density; DSP via Rankin — conceptual cite for the rHF-blind framing.
- [CB26] Brendan Chambers-Bourgeois. The root hermite factor is blind: Structural differences in lattice reduction revealed by the Li–Nguyen rankin profile. Zenodo, 2026. Companion paper (paper 1). Zenodo is its only release — no ePrint or other pre/post-publication. URL: <https://zenodo.org/records/20518060>, doi:10.5281/zenodo.20518060.
- [DEP23] Léo Ducas, Thomas Espitau, and Eamonn W. Postlethwaite. Finding short integer solutions when the modulus is small. In *CRYPTO 2023; IACR ePrint 2023/1125*, 2023. Z-shape q-ary geometry; Mitaka forgery at $\beta \approx 40$ — matches the $\beta = 40$ regime used here.
- [DvW21] Léo Ducas and Wessel van Woerden. NTRU fatigue: How stretched is overstretched? In *ASIACRYPT 2021*, 2021. Defines the circulant NTRU construction and the fatigue point $q_{\text{fat}} \approx 0.004 n^{2.484}$; ZGSA Z-shape profile.
- [J⁺23] Huiwen Jia et al. A new self-dual BKZ algorithm based on lattice sieving. In *FCS 2023, Springer CCIS 1992*, 2023. Self-dual BKZ with a G6K sieve oracle — closest precedent to the cross-engine self-dual construction here.
- [LN20] Jianwei Li and Phong Q. Nguyen. A complete analysis of the BKZ lattice reduction algorithm. Cryptology ePrint Archive, Report 2020/1237, 2020. URL: <https://eprint.iacr.org/2020/1237>.
- [MW16] Daniele Micciancio and Michael Walter. Practical, predictable lattice basis reduction. In Marc Fischlin and Jean-Sébastien Coron, editors, *EUROCRYPT 2016, Part I*, volume 9665 of *LNCS*, pages 820–849. Springer, Berlin, Heidelberg, May 2016. doi:10.1007/978-3-662-49890-3_31.
- [Row24] Rowell. *Theoretical, Practical, and Engineering Aspects of Lattice Cryptanalysis and TLS Attestation*. PhD thesis, Royal Holloway, University of London, 2024. Ch.3: approximate-enum oracles in BKZ; SD-BKZ achieves the same RHF as BKZ under a non-exact oracle (RHF-blind, independent).
- [Whi26] Trill White. Optimisation in BKZ: An analysis of dynamic block size. AMSI Summer Research Scholarships 2025–26, Deakin University, 2026. Dynamic block-size strategy within the Li–Nguyen framework. URL: https://srs.amsi.org.au/wp-content/uploads/sites/92/2026/03/white_trill_srs-report.pdf.